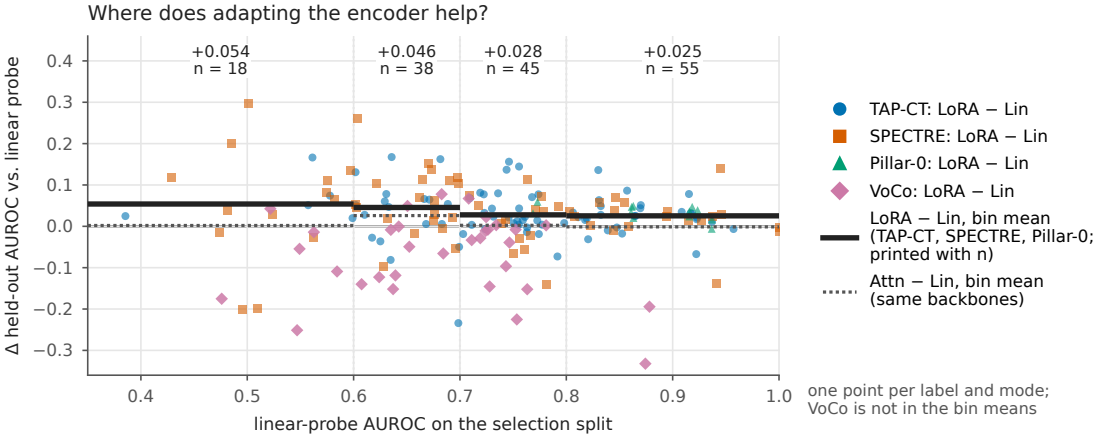
Where does adapting the encoder help?
+0.054
n = 18
+0.046
n = 38
+0.028
n = 45
+0.025
n = 55
Δ held-out AUROC vs. linear probe
linear-probe AUROC on the selection split
TAP-CT: LoRA − Lin
SPECTRE: LoRA − Lin
Pillar-0: LoRA − Lin
VoCo: LoRA − Lin
LoRA − Lin, bin mean
(TAP-CT, SPECTRE, Pillar-0;
printed with n)
Attn − Lin, bin mean
(same backbones)
one point per label and mode;
VoCo is not in the bin means