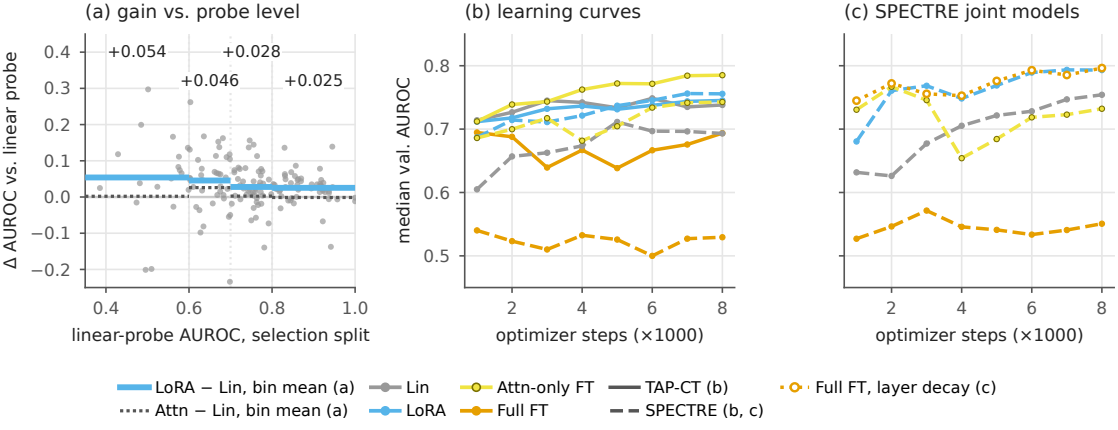
(a) gain vs. probe level
+0.054
+0.028
+0.046
+0.025
Δ AUROC vs. linear probe
linear-probe AUROC, selection split
(b) learning curves
median val. AUROC
optimizer steps (×1000)
(c) SPECTRE joint models
optimizer steps (×1000)
LoRA − Lin, bin mean (a)
Attn − Lin, bin mean (a)
Lin
LoRA
Attn-only FT
Full FT
TAP-CT (b)
SPECTRE (b, c)
Full FT, layer decay (c)