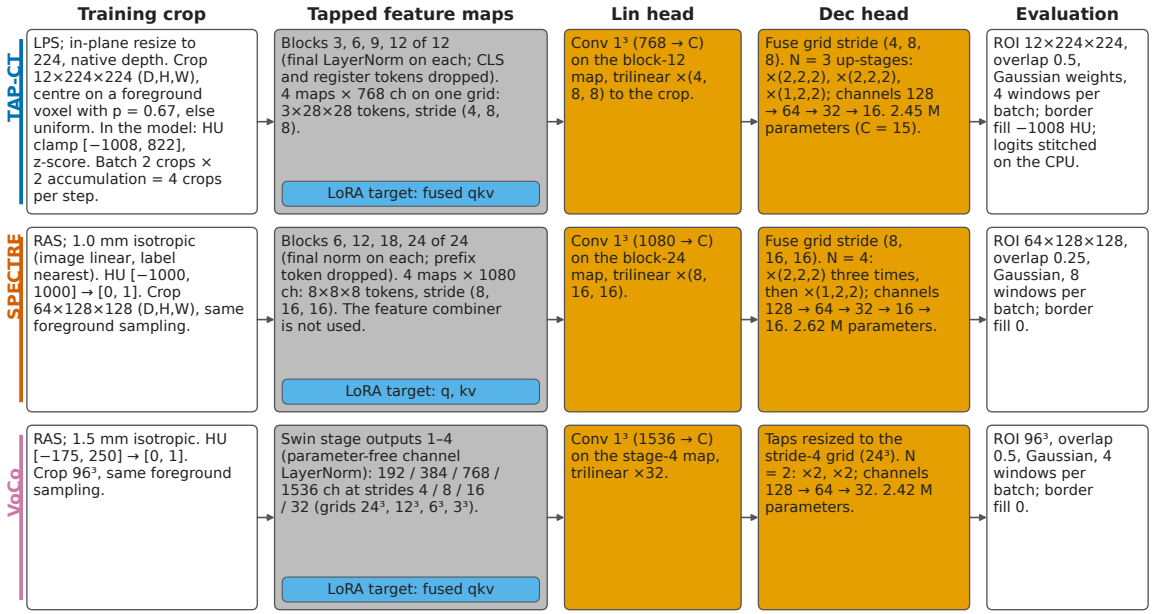

Dec head, layer by layer (N = 3 / 4 / 2 for TAP-CT / SPECTRE / VoCo; no skip connections from the input volume)

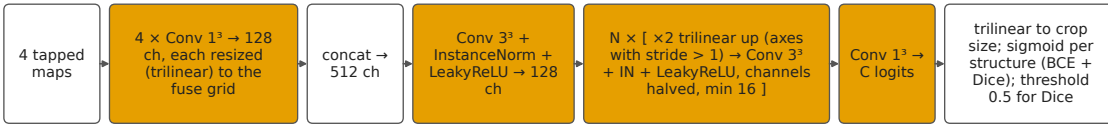

Lin head = Conv 1³ on the last tap + trilinear upsampling to the crop. C = number of structures (1 per-task, K joint). Encoder: frozen under Lin/Dec, adapted or trained under the other methods.

frozen (pretrained weights) low-rank adapter, trained trained in full